

Datalog

Datalog

- A nonprocedural language based on Prolog
 - Describe what instead of how: specifying the information desired without giving a specific procedure of obtaining that information
 - Resemble the syntax of Prolog
- A purely declarative manner
 - Simplify writing simple queries
 - Make query optimization easier


Basic Example

- Define a view relation $v1$ containing account numbers and balances for accounts at the Perryridge branch with a balance of over \$700
 - $v1(A, B) :- account(A, \text{"Perryridge"}, B), B > 700$
 - **for all** A, B
 - if** $(A, \text{"Perryridge"}, B) \in account$ **and** $B > 700$
 - then** $(A, B) \in v1$
- A Datalog program consists of a set of rules

Evaluation of a Datalog Program

- $v1(A, B) :- \text{account}(A, \text{"Perryridge"}, B), B > 700$

<i>account-number</i>	<i>branch-name</i>	<i>balance</i>
A-101	Downtown	500
A-215	Mianus	700
A-102	Perryridge	400
A-305	Round Hill	350
A-201	Perryridge	900
A-222	Redwood	700
A-217	Perryridge	750



<i>account-number</i>	<i>balance</i>
A-201	900
A-217	750

Retrieving Tuples

- Retrieve the balance of account number “A-217” in the view relation $v1$

$? v1(\text{“A-217”}, B)$

– Answer: (A-217, 750)

- Find account number and balance of all accounts in $v1$ that have a balance greater than 800

$? v1(A, B), B > 800$

– Answer: (A-201, 900)

A Program of Multiple Rules

- The interest rates for accounts

interest-rate(A, 5) :- account(A, N, B), B < 10000

interest-rate(A, 6) :- account(A, N, B), B >= 10000

- The set of tuples in a view relation is defined as the union of all the sets of tuples defined by the rules for the view relation

Negation

- Define a view relation c that contains the names of all customers who have a deposit but no loan at the bank

$c(N) :- \text{depositor}(N, A), \text{not } \text{is-borrower}(N).$
 $\text{is-borrower}(N) :- \text{borrower}(N, L)$

- Using **not** $\text{borrower}(N, L)$ in the first rule results in a different meaning, namely there is some loan L for which N is not a borrower
 - To prevent such confusion, we require all variables in negated “predicate” to also be present in non-negated predicates

Syntax of Datalog Rules

- Positive literal: $p(t_1, t_2 \dots, t_n)$
 - p is the name of a relation with n attributes
 - Each t_i is either a constant or variable
 - Example: `account(A, "Perryridge", B)`
- Negative literal: **not** $p(t_1, t_2 \dots, t_n)$
- Comparison and arithmetic are treated as positive predicates
 - $X > Y$ is treated as a predicate $>(X, Y)$
 - $A = B + C$ is treated as $+(B, C, A)$

Fact and Rules

- Fact $p(v_1, v_2, \dots, v_n)$
 - Tuple (v_1, v_2, \dots, v_n) is in relation p
- Rules: $p(t_1, t_2, \dots, t_n) \text{ :- } L_1, L_2, \dots, L_m.$

$\underbrace{\hspace{10em}}$
head

$\underbrace{\hspace{10em}}$
body

 - Each of the L_i 's is a literal
 - Head – the literal $p(t_1, t_2, \dots, t_n)$
 - Body – the rest of the literals
- A Datalog program is a set of rules

An Example Datalog Program

- Define interest on Perryridge accounts

```
interest(A, I) :- account(A, "Perryridge", B),  
                    interest-rate(A, R), I=B*R/100.  
interest-rate(A, 5) :- account(A, N, B), B<10000.  
interest-rate(A, 6) :- account(A, N, B), B>=10000.
```

Dependency of View Relations

- View relation v_1 depends directly on v_2 if v_2 is used in the expression defining v_1
 - Relation interest depends directly on relations interest-rate and account
- View relation v_1 depends indirectly on v_2 if there is a sequence of intermediate relations $v_1=i_1, \dots, i_n=v_2$ such that v_j depends directly on v_{j+1} for $1 \leq j < n$
 - Relation interest depends indirectly on relation account
- View relation v_1 depends on v_2 if v_1 depends directly or indirectly on v_2

Recursive Relation

- A view relation v is recursive if it depends on itself, otherwise, it is nonrecursive
- An example – defining the relation employment

$\text{empl}(X, Y) \text{ :- manager}(X, Y).$

$\text{empl}(X, Y) \text{ :- manager}(X, Z), \text{empl}(Z, Y)$

Semantics of Nonrecursive Datalog

- A ground instantiation of a rule (or simply instantiation) is the result of replacing each variable in the rule by some constant
 - Rule: $v1(A,B) :- \text{account}(A, \text{"Perryridge"}, B), B > 700.$
 - An instantiation:
 $v1(\text{"A-217"}, 750) :- \text{account}(\text{"A-217"}, \text{"Perryridge"}, 750), 750 > 700.$
- The body of rule instantiation R' is satisfied in a set of facts (database instance) I if
 - For each positive literal $q_i(v_{i,1}, \dots, v_{i,n_i})$ in the body of R' , I contains the fact $q_i(v_{i,1}, \dots, v_{i,n_i})$; and
 - For each negative literal **not** $q_j(v_{j,1}, \dots, v_{j,n_j})$ in the body of R' , I does not contain the fact $q_j(v_{j,1}, \dots, v_{j,n_j})$

Inferring Facts

- The set of facts that can be inferred from a given set of facts I using rule R as: $infer(R, I) = \{p(t_1, \dots, t_n) \mid \text{there is a ground instantiation } R' \text{ of } R \text{ where } p(t_1, \dots, t_n) \text{ is the head of } R', \text{ and the body of } R' \text{ is satisfied in } I\}$
- Given a set of rules $\mathfrak{R} = \{R_1, R_2, \dots, R_n\}$, define $infer(\mathfrak{R}, I) = infer(R_1, I) \cup infer(R_2, I) \cup \dots \cup infer(R_n, I)$

Example

- Rule: $v1(A,B) :- \text{account}(A, \text{"Perryridge"}, B), B > 700$

<i>account-number</i>	<i>branch-name</i>	<i>balance</i>
A-101	Downtown	500
A-215	Mianus	700
A-102	Perryridge	400
A-305	Round Hill	350
A-201	Perryridge	900
A-222	Redwood	700
A-217	Perryridge	750

A set of facts I


infer(R, I)

<i>account-number</i>	<i>balance</i>
A-201	900
A-217	750

Layer the View Relations

- Program

```
interest(A, I) :- perryridge-account(A, B),  
                  interest-rate(A, R), I = B * R/100.  
perryridge-account(A, B) :- account(A, "Perryridge", B).  
interest-rate(A, 5) :- account(N, A, B), B < 10000.  
interest-rate(A, 6) :- account(N, A, B), B >= 10000.
```



Layers

- A relation is in layer 1 if all relations used in the bodies of rules defining it are stored in the database
- A relation is in layer 2 if all relations used in the bodies of rules defining it are either stored in the database, or are in layer 1
- A relation p is in layer $i + 1$ if
 - It is not in layers 1, 2, ..., i
 - All relations used in the bodies of rules defining a p are either stored in the database, or are in layers 1, 2, ..., i

Semantics of a Program

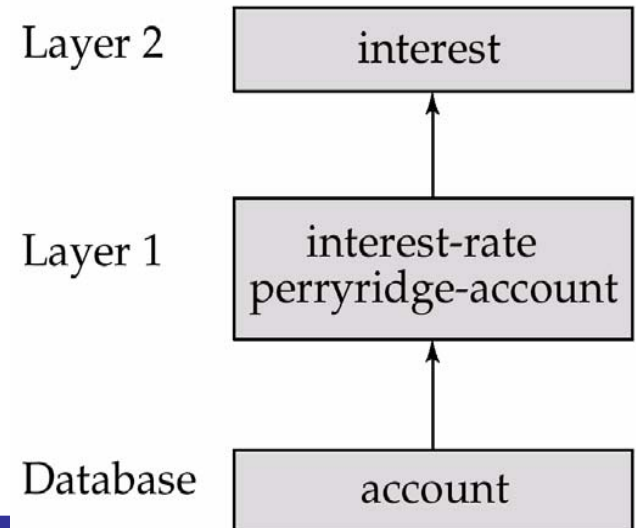
- Let the layers in a given program be $1, 2, \dots, n$.
Let \mathcal{R}_i denote the set of all rules defining view relations in layer i
- Define I_0 = the set of facts stored in the database
- Recursively define $I_{i+1} = I_i \cup \text{infer}(\mathcal{R}_{i+1}, I_i)$
- The set of facts in the view relations defined by the program (also called the semantics of the program) is given by the set of facts I_n corresponding to the highest layer n

Example

- Program

```
interest(A, I) :- perryridge-account(A, B),  
                  interest-rate(A, R), I = B * R/100.  
perryridge-account(A, B) :- account(A, "Perryridge", B).  
interest-rate(A, 5) :- account(N, A, B), B < 10000.  
interest-rate(A, 6) :- account(N, A, B), B >= 10000.
```

- I_0 : account
- I_1 : account, interest-rate
- I_2 : account, interest-rate, interest



Safety

- Unsafe rules – lead to infinite answers
 - $gt(X, Y) :- X > Y$
 - $not\text{-}in\text{-}loan(B, L) :- \textbf{not } loan(B, L)$
 - $P(A) :- q(B)$
- Safety conditions
 - Every variable that appears in the head of the rule also appears in a non-arithmetic positive literal in the body of the rule
 - Every variable appearing in a negative literal in the body of the rule also appears in some positive literal in the body of the rule
- If a nonrecursive Datalog program satisfies the safety conditions, then all the view relations defined in the program are finite

Relational Operations

- Project out attribute *account-name* from account.

$query(A) :- account(A, N, B).$

- Cartesian product of relations r_1 and r_2 .

$query(X_1, X_2, \dots, X_n, Y_1, Y_1, Y_2, \dots, Y_m) :-$
 $r_1(X_1, X_2, \dots, X_n), r_2(Y_1, Y_2, \dots, Y_m).$

- Union of relations r_1 and r_2 .

$query(X_1, X_2, \dots, X_n) :- r_1(X_1, X_2, \dots, X_n),$
 $query(X_1, X_2, \dots, X_n) :- r_2(X_1, X_2, \dots, X_n),$

- Set difference of r_1 and r_2 .

$query(X_1, X_2, \dots, X_n) :- r_1(X_1, X_2, \dots, X_n),$
not $r_2(X_1, X_2, \dots, X_n)$

Recursion

Relation schema $\text{manager}(\text{employee}, \text{manager})$

$\text{empl-jones}(X) \text{ :- } \text{manager}(X, \text{Jones}).$

$\text{empl-jones}(X) \text{ :- } \text{manager}(X, Y), \text{empl-jones}(Y).$

<i>employee-name</i>	<i>manager-name</i>
Alon	Barinsky
Barinsky	Estovar
Corbin	Duarte
Duarte	Jones
Estovar	Jones
Jones	Klinger
Rensal	Klinger

Iteration number	Tuples in <i>empl-jones</i>
0	
1	(Duarte), (Estovar)
2	(Duarte), (Estovar), (Barinsky), (Corbin)
3	(Duarte), (Estovar), (Barinsky), (Corbin), (Alon)
4	(Duarte), (Estovar), (Barinsky), (Corbin), (Alon)

Datalog Fixpoint

- The view relations of a recursive program containing a set of rules \mathcal{R} are defined to contain exactly the set of facts I computed by the iterative procedure *Datalog-Fixpoint*

procedure Datalog-Fixpoint

I = set of facts in the database

repeat

$Old_I = I$

$I = I \cup infer(\mathcal{R}, I)$

until $I = Old_I$

- At the end of the procedure, $infer(\mathcal{R}, I) \subseteq I$
 - $infer(\mathcal{R}, I) = I$ if we consider the database to be a set of facts that are part of the program
- I is called a fixed point of the program

Semantics of Recursion

- Fixpoint
 - Fixpoint is unique
- Transitive closure of a relation
 - $empl(X, Y) :- manager(X, Y).$
 $empl(X, Y) :- manager(X, Z), empl(Z, Y)$
- Another way
 - $empl(X, Y) :- manager(X, Y).$
 $empl(X, Y) :- empl(X, Z), manager(Z, Y).$
- Cannot use negation

The Power of Recursion

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration
- Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins
- Programs satisfy the safety condition will terminate
 - `number(0). number(A) :- number(B), A=B+1.`
 - Some programs not satisfying the safety condition do terminate

Monotonicity

- A view V is said to be monotonic if given any two sets of facts I_1 and I_2 such that $I_1 \subseteq I_2$, then $E_V(I_1) \subseteq E_V(I_2)$, where E_V is the expression used to define V
- A set of rules R is said to be monotonic if
$$I_1 \subseteq I_2 \text{ implies } \text{infer}(R, I_1) \subseteq \text{infer}(R, I_2),$$
- Relational algebra views defined using only the operations: Π , σ , \times , \cup , \cap , and ρ are monotonic
 - Relational algebra views defined using “–” may not be monotonic.
- Datalog programs without negation are monotonic, but Datalog programs with negation may not be monotonic
- Monotonic expressions can use the fixpoint technique

Summary

- Datalog: a prolog-like query language
- Using Datalog to write queries
- Semantics of Datalog programs

To-Do-List

- Examine the example queries in the relational algebra section, which ones can be rewritten in Datalog?