

# Real-World Datomic: An Experience Report



# Me

- Craig Andera
- @craigandera
- craig@cognitect.com
- <http://cognitect.com/podcast>

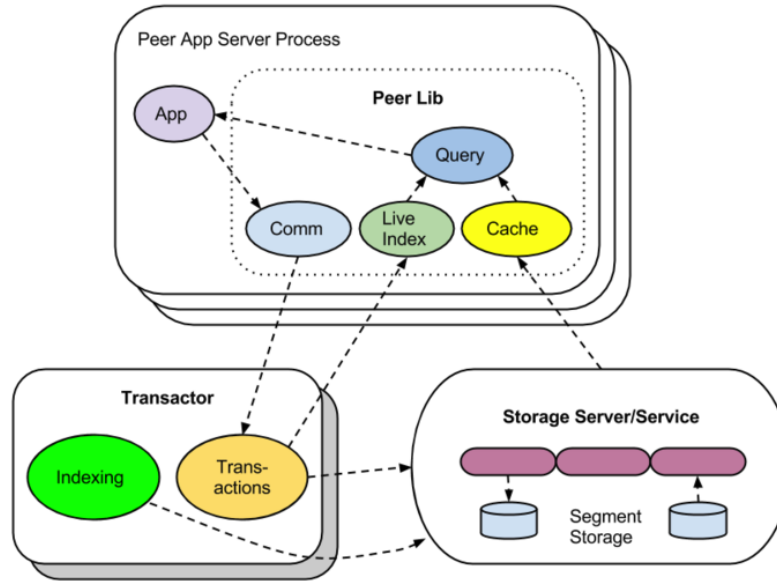
# Room Key

- <http://roomkey.com>
- Employer of Clojure Programmers
- <http://is.gd/rkpodcast>

# The Problem

- Provide a summary of sales funnel events for Room Key
- Users move through the website
  - impression -> hotel details -> lead -> booking
- Want per-day event counts by hotel & event type
- $30 \text{ days} * \sim 20 \text{ events/sec} * \sim 5 \text{ hotels/event}$ 
  - =>  $\sim 250\text{M}$  hotel-events

# Datomic's Architecture



# Datomic's Data Model

- All facts stored as **datoms**
- A datom is EAVT tuple
  - **Entity** - The thing we're modeling
  - **Attribute** - Which aspect we're describing
  - **Value** - The quantity or quality
  - **Time/Tx** - When we learned this fact
- *Yesterday, as part of a transaction that happened at 3:02PM, I learned that Craig likes pizza*

# Attributes

- Have a type, an ident, and a cardinality
  - Optionally other stuff, e.g. uniqueness
- Types include the usual scalars: strings, numbers, etc.
- Also include **refs**
  - Value is an entity ID
  - Allows datoms to form graphs
- Set of defined attributes constitutes database schema

## Capturing Events - Schema

Attribute	Type	Qualifiers
<code>:roomkey.hotel/id</code>	string	unique
<code>:roomkey.event/type</code>	ref	single
<code>:roomkey.event/hotels</code>	ref	many
<code>:roomkey.event/time</code>	inst	single



# Capturing Events - Example

Impression (EID 123)	
:db/ident	:impression

Hotel details (EID 234)	
:db/ident	:details

Lead (EID 345)	
:db/ident	:lead

Booking (EID 456)	
:db/ident	:booking

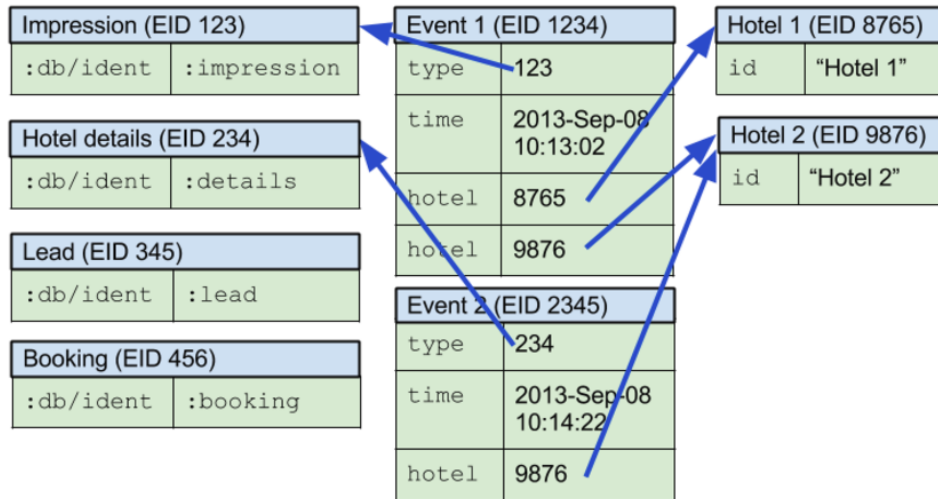
Event 1 (EID 1234)	
type	123
time	2013-Sep-08 10:13:02
hotel	8765
hotel	9876

Event 2 (EID 2345)	
type	234
time	2013-Sep-08 10:14:22
hotel	9876

Hotel 1 (EID 8765)	
id	"Hotel 1"

Hotel 2 (EID 9876)	
id	"Hotel 2"

# Capturing Events - Refs



# Storing and Retrieving Data in Datomic

- One or more EAV tuples submitted to the transactor
- All are given the same  $t$  value
- Datoms written transactionally to the log
- All datoms also stored redundantly in the **indexes**

# Datomic Indexes

- Every datom stored in two or more **indexes**
- Index is (logically) a sorted set of datoms
- Indexes named by their sort order
- EAVT, AEVT, AVET, and VAET

# EAVT and AEVT

- All datoms stored in these
- EAVT
  - Efficient access to all attributes of an entity
- AEVT
  - Efficient access to all values of an attribute

## Capturing Events - EAVT

E	A	V	T
<Event 1>	type	<impression>	1111
<Event 1>	time	10:13:02	1111
<Event 1>	hotel	<Hotel 1>	1111
<Event 1>	hotel	<Hotel 2>	1111
<Event 2>	type	<details>	2222
<Event 2>	time	10:14:22	2222
<Event 2>	hotel	<Hotel 2>	2222

# Capturing Events - EAVT

E	A	V	T
1234	26	123	1111
1234	27	10:13:02	1111
1234	28	8765	1111
1234	28	9876	1111
2345	26	234	2222
2345	27	10:14:22	2222
2345	28	9876	2222

## Capturing Events - AEVT

E	A	V	T
<Event 1>	type	<impression>	1111
<Event 2>	type	<details>	2222
<Event 1>	time	10:13:02	1111
<Event 2>	time	10:14:22	2222
<Event 2>	hotel	<Hotel 1>	1111
<Event 1>	hotel	<Hotel 2>	1111
<Event 2>	hotel	<Hotel 2>	2222



# AVET and VAET

- AVET
  - Only stores datoms with attribute marked :**db/index**
  - Efficient lookup of entities by attribute/value pairing
- VAET
  - Only stores datoms with attribute of type **ref**
  - Allows efficient navigation of relationships in reverse
  - Also called the **reverse index**

# Capturing Events - AVET

E	A	V	T
<Hotel 1>	id	"Hotel 1"	888
<Hotel 2>	id	"Hotel 2"	999


## Capturing Events - VAET

E	A	V	T
<Event 1>	type	<impression>	1111
<Event 2>	type	<details>	2222
<Event 1>	hotel	<Hotel 1>	1111
<Event 1>	hotel	<Hotel 2>	1111
<Event 2>	hotel	<Hotel 2>	2222


# Index Storage and Retrieval

- Datoms are stored in compressed chunks called **segments**
- Segments are stored as a tree
- Segments are immutable once written
- Each transaction conceptually creates a new root
- Persistent indexes built less frequently
- Peers pull segments from the appropriate index as needed


# Index Trees



# Indexing - Append



# Indexing - Highly Nonlocal



## Trying to Solve the Problem - Query

```
[ :find ?event-type ?hotel-id ?day (count ?event)
  :in $  [?hotel-id ...] ?since
  :where
    [?hotel-eid :roomkey.hotel/id ?hotel-id]
    [?event :roomkey.event/hotel ?hotel-eid]
    [?event :roomkey.event/time ?time]
    [(<= ?since ?time)]
    [(roomkey.time/midnight ?time) ?day]
    [?event :roomkey.event/type ?event-type]]
```



# Why This Didn't Work

- Too much data, not enough time: ~250M datoms, 100ms
- No query order eliminates enough data

# Anatomy of an Entity ID

- Entity IDs have two parts
  - High bits are the **partition** ID
  - Low bits are a timelike component
- Partition ID user-assigned at entity creation

# A Change In Approach

- A big, giant cache
- Stop recording time explicitly
- Record events using a dedicated partition
- Use **seek-datoms**
  - Returns `java.lang.Iterable` over raw index data
  - Fast, lazy, reduceable
- **entid-at**
  - Given a time and a partition, get an EID

# The Code

```
(defn datoms-between
  "Returns a reducible collection of datoms created
  between the start and end dates in a single partition."
  [db partition start end]
  (let [start-e (d/entid-at db
                             partition
                             start)
        end-e   (d/entid-at db
                             partition
                             end)]
    (->> (d/seek-datoms db :eavt start-e)
          (r/take-while #(< (:e %) end-e)))))
```

# Why You Shouldn't Always Do This

- Datalog is a big deal
  - Joins!
  - Query as data
  - Future improvements to query
- Segment caching: hot data often in-memory

# Colophon

- Appreciations
  - Room Key
  - Tim Ewald
  - Rich Hickey
  - Bobby Calderwood
  - Stuart Halloway
- Typography
  - Carrois Gothic

**Questions?**

**Thanks!**